

The Cost of Saying Hello:

An Empirical Study of Handshake Overhead in AI Agent Communication

Protocols

Pranav Varshney
University of Michigan
pvarsh@umich.edu

Abstract

Every interaction between AI agents begins with a handshake. The initiating agent must locate its counterpart, learn what it can do, and establish a session before the first useful message can be sent. Surveys of agent communication protocols name handshake overhead as an efficiency criterion, yet none report measurements. This paper presents the first empirical, phase-decomposed benchmark of the handshake in three major agent protocols: MCP (in both its 2025 *initialize* and 2026 *discover* eras), A2A, and ACP. Using each protocol’s official SDK, we instrument every request for round trips, per-phase latency, wire bytes, and token cost across cold-start, warm-repeat, capability-mismatch, and degraded-metadata scenarios ($N=50$ per cell, on loopback and under an emulated 50 ms RTT). Round-trip count dominates handshake cost. A legacy MCP handshake takes three round trips and 230 ms at 50 ms RTT; A2A needs one round trip and 61.0 ms. Round trips are not the whole bill: when an LLM reads the counterpart’s metadata, handshake cost scales with the counterpart’s capability inventory at 24–107 tokens per capability, so a 200-capability counterpart spends 11% of a 200k-token context window before the agent reads its own task. MCP’s own 2026 revision cuts its handshake from three round trips to two and, with era pinning, to zero on warm reconnects, protocol evolution that corroborates the cost we measure. We release the full harness, raw data, and analysis at <https://github.com/pvarshh/handshake-benchmark>.

Keywords: agent communication protocols, Model Context Protocol, A2A, handshake overhead, connection establishment, benchmarking

1 Introduction

Software agents built on large language models are beginning to interact with other agents as a matter of course. A coding assistant calls a tool server; an orchestrator hands a subtask to a specialist; a consumer agent queries a merchant agent. Each of these interactions is preceded by a connection-establishment exchange in which the initiating agent discovers its counterpart, learns its capabilities, and establishes a session. We call this exchange the *handshake*, by analogy with transport-layer connection setup. The handshake is pure overhead. No useful work happens until it completes, and in open agent ecosystems, where counterparts are strangers, capabilities are heterogeneous, and interactions are often short, its cost is paid frequently and amortized poorly.

The systems community has a long tradition of measuring exactly this kind of tax at lower layers. The TCP three-way handshake, the TLS key exchange [12, 18, 19], and QUIC’s 0-RTT resumption [11] were all designed, measured, and redesigned around one insight: connection establishment is a first-order cost of short interactions. The emerging stack of AI agent communication protocols has not yet received the same treatment. The most comprehensive survey of the space names “handshake overhead” as an explicit efficiency criterion for comparing protocols, yet reports no measurements [22]. Recent benchmarks measure agent protocols end-to-end at the task level [4, 5], but none isolates the connection-establishment phase itself.

This paper fills that gap. We benchmark the handshake, and only the handshake, of three major agent communication protocols head-to-head, using their official SDKs: the Model Context Proto-

col (MCP) [3, 17], the Agent2Agent protocol (A2A) [1], and the Agent Communication Protocol (ACP) [10]. ACP merged into A2A in August 2025 and is frozen; we include it because its deployments persist and because its one-round-trip REST design is a distinct point in the handshake design space, which a comparison of only the two live protocols would miss. For MCP we measure two protocol eras separately. The *legacy* handshake (specification revisions 2024-11 through 2025-11) is what nearly all deployed MCP software used at the time of writing; the *modern* handshake, introduced in revision 2026-07-28, replaces the three-message initialization with a single cacheable probe. This within-protocol comparison turns a recent design change into a natural experiment on the cost of handshake round trips.

Our contributions are: (C1) a formal decomposition of the agent handshake into discovery, capability-exchange, session, and readiness phases, with a precise per-protocol definition of the *task-ready* state (Section 3); (C2) to our knowledge the first cross-protocol empirical benchmark of handshake overhead, covering round trips, phase-decomposed latency, wire bytes, capability-metadata token cost, and failure behavior across four scenario families and two network conditions, with $N=50$ runs per cell (Sections 4–5); (C3) an open reproducibility package (harness, minimal counterpart agents, raw per-run logs, and analysis scripts) with one-command reproduction; and (C4) design implications for protocol and agent-framework authors, grounded in where the measured overhead actually concentrates (Section 6).

Our headline findings are simple. First, the moment a real network is involved, round-trip count dominates handshake cost; payload size and compute do not. Under a 50 ms RTT, the three-round-trip legacy MCP handshake costs 230 ms to task-ready (median), while A2A’s single card fetch costs 61.0 ms, a factor of 3.8. Second, protocols differ sharply in what a warm reconnection may skip. A2A, ACP, and modern MCP can reach task-ready with zero round trips using spec-sanctioned caching (modern MCP when pinned to a cached protocol era; its default policy still revalidates in one), while legacy MCP repeats its full handshake on every connection. Third, failure behavior diverges more than success behavior. All three protocols detect malformed metadata quickly, but only MCP validates version claims during the handshake itself. A2A, by design, defers version checking to a per-request header; its client completed tasks against a counterpart whose card claims protocol version 99.0.0, and a real incompatibility would surface only at the first task message.

2 Background and Related Work

2.1 The protocols under study

MCP. The Model Context Protocol, introduced by Anthropic in November 2024 [3] and now governed as a Linux Foundation project, connects a client (typically an LLM application) to servers exposing tools, resources, and prompts. It is a JSON-RPC protocol; we use its streamable HTTP transport. Through specification revision 2025-11-25 the handshake comprises an `initialize` request/response carrying protocol version and capability flags, an `initialized` notification, and, for a client that needs the tool inventory, a `tools/list` request [16]. Revision 2026-07-28 introduces a second era. A single `server/discover` probe returns versions, capabilities, and instructions in one exchange, and responses to discovery and listing requests carry explicit freshness hints (`ttlMs`, `cacheScope`; SEP-2549) that make them client-cacheable across sessions [17]. We measure both eras.

A2A. The Agent2Agent protocol, announced by Google in April 2025 and donated to the Linux Foundation, targets peer-to-peer task delegation between opaque agents. Discovery is document-centric: an agent publishes an *Agent Card* at a well-known URI (`/.well-known/agent-card.json`) describing its identity, supported interfaces, and skills, and task exchange then proceeds over a JSON-RPC, REST, or gRPC binding [1]. We benchmark specification v1.0.0 with the JSON-RPC binding.

ACP. The Agent Communication Protocol, from IBM’s BeeAI project, is REST-first. A server exposes `GET /agents`, returning manifests for the agents it hosts, and runs are created with `POST`

/runs. In August 2025 the ACP project merged into A2A under the Linux Foundation and its repository was archived [13]; the specification and SDK remain published and deployed [10]. We include ACP because deployments persist and because its one-round-trip REST design occupies a distinct point in the handshake design space. We treat it throughout as a frozen protocol.

2.2 Related work

Surveys of the agent-protocol landscape compare these protocols qualitatively. Yang et al. [22] propose evaluation dimensions, efficiency and handshake overhead explicitly among them, but decline to define a benchmark. Ehtesham et al. [6] compare MCP, ACP, A2A, and ANP architecturally. Security analyses have compared the protocols’ threat surfaces [2, 9, 14]. Performance has received far less attention.

Two recent empirical efforts are the closest relatives of this work. ProtocolBench [5] benchmarks A2A, ACP, ANP, and Agora on task success, end-to-end latency, message overhead, and robustness across multi-agent workloads; MCP is explicitly out of scope, and its timeline instrumentation (send, queue, service, first-token) does not isolate connection establishment. Dobrovolskyi [4] compares an MCP-based single-agent architecture against A2A multi-agent architectures end-to-end over 450 task executions. That study attributes part of the difference to setup costs, but it infers them rather than instrumenting them. Neither work decomposes or measures the handshake itself. Our benchmark is complementary to both. Task-level benchmarks tell you what a protocol costs once agents are talking, and ours tells you what it costs to start talking.

The Agora protocol [15] occupies a different corner of the design space: agents negotiate protocols in natural language, trading tokens for flexibility. We do not benchmark live LLM-mediated negotiation here (Section 7), but Agora’s own measurements of negotiation cost motivate our token-cost metric.

The closest living analogue is service discovery in microservices, where the same costs are paid at a different layer. Envoy’s xDS protocol streams endpoint and cluster configuration to sidecars so that a caller never pays discovery on the request path [7], and HTTP/2 connection reuse and multiplexing exist precisely to keep per-exchange setup off that path [21]. Both are architectural answers to the question we measure: how do you stop paying to find out who you are talking to? The methodological precedent is also close. Zhu et al. [23] decompose service-mesh sidecar overhead into its constituent components and predict it from configuration, which is the same compositional move we make for the handshake. The difference is what the two layers can assume. A service mesh operates over a known, centrally configured fleet, so discovery state can be pushed ahead of time; open agent ecosystems have no such control plane, and the counterpart is a stranger whose capabilities are discovered on first contact. That is why the costs we measure are paid per pair rather than amortized fleet-wide.

Classical connection-establishment measurement is a direct methodological ancestor. Naylor et al. [18] quantified the latency cost of the TLS handshake at Internet scale, Lee et al. [12] measured TLS 1.3’s handshake-time reduction in deployment, and QUIC folded transport and cryptographic setup into one round trip with 0-RTT resumption [11]. Pre-LLM multi-agent systems paid setup costs too. Contract Net’s announce–bid–award cycle is a capability-discovery round trip [20], and FIPA-ACL interaction protocols begin with capability registration and service discovery [8]. What is genuinely new in the LLM era is that handshake artifacts are often consumed by a language model. An agent card or tool list may be placed into an LLM’s context window, so handshake content has a per-token *inference* cost that classical protocol analysis never had to account for. Byte counts and token counts are correlated but priced very differently; we therefore measure both.

3 Defining the Agent Handshake

We define the *handshake* as all communication from first contact until the initiating agent can send its first substantive task message with justified confidence that the counterpart can handle it. We de-

Table 1: Task-ready definitions and observed cold-start wire sequences. Request and response sizes are HTTP body bytes from instrumented runs.

Protocol	Task-ready condition	Observed wire sequence (cold)
MCP (legacy)	<code>initialize</code> response received; <code>initialized</code> notification sent; <code>tools/list</code> parsed; <code>echo</code> present	<code>POST initialize</code> (163 B→379 B); <code>POST initialized</code> (54 B, 202); <code>POST tools/list</code> (68 B→399 B); 3 round trips
MCP (modern)	<code>server/discover</code> result adopted; <code>tools/list</code> parsed; <code>echo</code> present	<code>POST server/discover</code> (245 B→431 B); <code>POST tools/list</code> (240 B→522 B); 2 round trips
A2A	Agent Card fetched and parsed; skill <code>echo</code> present in card	<code>GET /.well-known/agent-card.json</code> (→443 B); 1 round trip
ACP	Manifest list fetched and parsed; agent <code>echo-agent</code> present	<code>GET /agents</code> (→507 B); 1 round trip

compose it into four phases: (a) *discovery*: locating the counterpart and retrieving its self-describing metadata; (b) *capability exchange*: learning what the counterpart can do, at whatever granularity the protocol exposes; (c) *auth/session*: establishing identity and any session state; and (d) *readiness*: any residual signaling required before the first task message is valid. A phase may be empty for a given protocol, and a single message may serve more than one phase; we attribute each wire exchange to the phase that constitutes its primary function.

The handshake ends at the *task-ready* state. Because “ready” is protocol-specific, we define it operationally per protocol (Table 1) against a fixed canonical capability: an `echo` operation that returns its input verbatim. The initiator is task-ready when it has positively confirmed, from protocol metadata, that the counterpart offers `echo`. Using a trivial capability keeps task execution cost negligible, so the handshake dominates. We also execute one `echo` call after each timed handshake, outside the timed region, to validate that the ready state is real.

Phase attribution follows the protocols’ semantics. MCP has no in-protocol discovery; the endpoint URL is assumed known, which is itself a design choice. Its legacy handshake therefore spends one exchange on session establishment (`initialize`, which also negotiates the protocol version and capability flags and yields the session identifier), one on readiness (`initialized`), and one on capability exchange (`tools/list`). Modern MCP’s `server/discover` performs version and capability negotiation without creating a session; we attribute it, and `tools/list`, to capability exchange. A2A and ACP fold discovery and capability exchange into a single metadata fetch, with capability confirmation a local parse. Neither establishes a session or requires auth in the anonymous configuration we benchmark.

We exclude from handshake measurement everything below the application layer (TCP and TLS setup, reported separately as a baseline in Section 5.6), task execution itself, and connection tear-down. We do count the TCP connections each handshake opens, since each one implies an additional transport round trip on a real network.

4 Methodology

4.1 Harness

The benchmark harness runs each protocol’s official Python SDK client against a minimal local counterpart agent built with the same project’s official server SDK. Counterparts are deliberately minimal, exposing static capability metadata plus the `echo` capability, so that measured costs are protocol floors rather than implementation artifacts. All three SDKs use `httpx` for HTTP, which gives us a uniform instrumentation point: a recording transport wraps every request with monotonic timestamps, body sizes, and reconstructed wire sizes, and tags it with a handshake phase. All traffic additionally traverses a local TCP relay that counts exact bytes and connections per run.

Network conditions are emulated in the relay. For the *50 ms RTT* condition it delays each forwarded chunk by 25 ms per direction, approximating a symmetric 50 ms round-trip link at the application layer. (We use a userspace relay rather than kernel-level emulation such as `tc netem` for portability; Section 4.5 discusses the difference.) The *localhost* condition uses the same relay with zero added delay, isolating protocol-stack overhead. For the cold-start scenario we additionally sweep the emulated RTT across 0, 25, 50, and 100 ms ($N=20$ runs per protocol per point, with fresh counterpart servers per cell) to measure how latency-to-ready scales with round-trip time (Figure 3).

Each run is fully cold on the client side: a fresh SDK client object, fresh HTTP connection pool, and fresh relay. Counterpart server processes persist across the runs of a cell, as a long-running service would, and five warmup runs precede each cell and are discarded. Timing uses a monotonic nanosecond clock. The handshake timer starts immediately before the client initiates first contact and stops when the task-ready condition is confirmed. Every run is logged as one JSON record containing the full request/response event list, and all statistics in this paper are computed from those records by the released analysis scripts; the numbers in the text are generated macros, not hand-typed values.

4.2 Scenarios

S1: cold start. Stranger agents. No prior contact, no cached metadata, full discovery to task-ready. This is the headline number.

S2: warm repeat. The same pair, subsequent connection. Each protocol is allowed exactly the reuse its specification and official SDK support. For MCP (legacy) that is nothing: response caching arrived only with the modern era, so the SDK re-runs the full three-round-trip handshake, tool listing included. For MCP (modern) we measure two policies. Under *auto*, the SDK’s default connect policy, the client re-probes `server/discover` and serves `tools/list` from the SEP-2549 response cache; under *pinned*, it additionally remembers the counterpart’s protocol era and adopts it directly, reaching task-ready entirely from cache. For A2A and ACP the client reuses its cached Agent Card or manifest. Our MCP counterpart advertises a one-hour public cache TTL, the configuration a latency-conscious deployment would choose; the protocol default is non-cacheable.

S3: capability mismatch. The initiator requires a `translate` capability the counterpart does not offer. We measure time-to-rejection, that is, how much of the handshake must be paid before the mismatch is discoverable.

S4: degraded conditions. (a) *Slow counterpart*: the counterpart delays every response by 2 s. (b) *Malformed metadata*: the counterpart returns its capability metadata (`initialize/discover` result, Agent Card, or manifest) truncated at 50%, yielding syntactically invalid JSON with valid HTTP framing. (c) *Version mismatch*: the counterpart advertises only a fictitious protocol version (2099-01-01 for MCP; 99.0.0 in the A2A Agent Card). ACP’s manifest carries no protocol-version field, so (c) is not testable for ACP, which is itself a finding. For each degraded cell we record outcome, time to failure or readiness, client retries, and the error surfaced to the caller.

4.3 Metrics and reporting

For every scenario–protocol–network cell we report the following: application-layer round trips to task-ready (request–response exchanges, including fire-and-forget notifications); wall-clock latency to task-ready, decomposed by phase; TCP connections opened; exact bytes on the wire in each direction (relay-counted, through task-ready); and, for degraded cells, retries and time-to-failure. We run $N=50$ measured runs per cell and report medians and 95th percentiles throughout, never bare means. Distributions were inspected for bimodality; none required special treatment after warmup exclusion. Every median carries a percentile bootstrap 95% confidence interval over 2,000 resamples, which separates two error sources that are easy to conflate. Within a cell, sampling precision is high: the median CI half-width is 0.16 ms and the largest is 2.0 ms, so a figure quoted to 0.1 ms is resolved by the data. Across cells, client-side setup time drifts by a few milliseconds over a long session (Section 4.5), and that systematic term, not sampling noise, is what limits between-protocol

comparison. We therefore quote medians at the precision the within-cell interval supports and state the cross-cell floor explicitly wherever we compare cells (Section 5.3).

For token cost we ask what the handshake costs when an LLM-driven initiator reads the counterpart’s capability metadata, the normal case when a model chooses tools or validates a counterpart. We take the exact metadata bodies observed on the wire, count tokens with the open o200k_base tokenizer, and price them at a mid-tier frontier model’s list input rate (Claude Sonnet 5, \$3.00 per million input tokens).¹ Anthropic’s own tokenizer counts roughly 15–20% more tokens on JSON-like content than o200k_base, so the dollar figures are conservative lower bounds. Byte counts are exact.

4.4 Environment and versions

All experiments ran on a single Apple-silicon machine (macOS, Darwin 25.5.0, arm64) on 2026-08-10, with client, relay, and counterpart on loopback. Python 3.12.13; mcp 2.0.0 (SDK released alongside spec revision 2026-07-28); a2a-sdk 1.1.2 (A2A spec v1.0.0, JSON-RPC binding); acp-sdk 1.0.3 (final release); httpx 0.28.1; uvicorn 0.34.3; fastapi 0.115.14. Exact pins and the one-command reproduction path are in the repository; the appendix lists per-protocol launch details.

4.5 Threats to validity

Toy counterparts. Our agents expose one capability. Real counterparts with large tool inventories would shift byte and token costs upward, though handshake latency is less affected, since inventory size changes payloads rather than round trips. Our numbers are protocol floors.

Userspace network emulation. The relay delays application-layer chunks, not packets. TCP connection establishment against the relay completes at loopback speed, so our 50 ms-RTT figures exclude transport setup entirely; on a real 50 ms link each TCP connection would add one additional RTT (plus one for TLS 1.3), which penalizes exactly the protocols that open more connections (Section 5.1). SSE-framed responses cross the relay in two chunks and therefore pay the per-chunk delay twice, which parallels, but does not perfectly model, per-packet kernel emulation.

Single SDK generation. We measure specific SDK versions at a specific date. SDK-added behaviors (connection handling, retries, caching) are part of what we measure, deliberately, because they are what deployed agents actually pay; results will drift as SDKs evolve.

Single machine, loopback. Absolute low-RTT latencies reflect one machine’s stack. The round-trip and byte counts, which drive the higher-RTT results, are machine-independent. Client-side setup time also drifts by a few milliseconds across cells within a long session, setting a ~5 ms floor on cross-cell latency comparisons (Section 5.3).

Token accounting. Token costs model a hypothetical LLM-in-the-loop reader with an open tokenizer at one vendor’s list price. They scale linearly in both assumptions and are stated as lower bounds.

5 Results

5.1 S1: cold start

Table 2 and Figure 1 give the cold-start picture. The protocols reach task-ready in three (MCP (legacy)), two (MCP (modern)), and one (A2A, ACP) application-layer round trips. On loopback, where a round trip costs microseconds, the differences are visible but small in absolute terms: medians of 6.1 ms, 4.4 ms, 2.8 ms, and 2.8 ms respectively. Client-side protocol-stack work dominates these numbers, not the wire.

Under the emulated 50 ms RTT the ordering is unchanged but the spread widens to exactly what the round-trip counts predict: 230 ms for MCP (legacy) (p95 236 ms), 122 ms for MCP (modern),

¹Anthropic pricing documentation, <https://platform.claude.com/docs/en/about-claude/pricing>, accessed 2026-08-11. Introductory pricing of \$2 per million input tokens is in effect through 2026-08-31; the standard \$3.00 rate takes effect 2026-09-01. We use the standard rate; all dollar figures scale linearly with the assumed price.

Table 2: Main results. Medians over $N=50$ runs per cell; RT = application-layer round trips to task-ready (or to rejection for S3); Conn = TCP connections opened; wire bytes are relay-counted through task-ready (localhost condition shown). p95 in adjacent columns.

Scenario	Protocol	RT	Conn	localhost (ms)		50 ms RTT (ms)		wire bytes	
				med	p95	med	p95	up	down
S1 cold start	MCP (legacy)	3	3	6.1	8.0	230	236	1147	1522
	MCP (modern)	2	1	4.4	5.2	122	125	1070	1205
	A2A	1	1	2.8	3.0	61.0	61.9	170	569
	ACP	1	1	2.8	3.2	62.1	63.1	149	633
S2 warm repeat	MCP (legacy)	3	3	6.1	6.5	230	234	1147	1522
	MCP (modern, auto)	1	1	3.6	4.2	63.5	65.4	540	557
	MCP (modern, pinned)	0	0	2.1	2.6	6.6	7.3	0	0
	A2A	0	0	2.0	2.7	2.0	2.7	0	0
	ACP	0	0	2.1	2.5	6.6	7.4	0	0
S3 mismatch	MCP (legacy)	3	3	6.2	6.9	231	235	1147	1522
	MCP (modern)	2	1	4.4	5.5	117	123	1070	1205
	A2A	1	1	2.9	4.3	62.2	65.3	170	569
	ACP	1	1	2.9	4.0	57.7	63.2	149	633

61.0 ms for A2A, and 62.1 ms for ACP. The legacy MCP handshake costs $3.8 \times$ A2A’s. Two details deserve note. First, MCP (legacy) exceeds the naive 3×50 ms prediction because two of its three responses are SSE-framed and cross the relay in two delayed chunks; measured against per-packet emulation the same effect appears as serialization of header and body packets. Second, in the legacy era the MCP client (mcp 2.0.0 defaults) opens a fresh TCP connection for every JSON-RPC exchange: 3 connections for the legacy handshake versus one each for modern MCP, A2A, and ACP. We verified this is client behavior rather than a harness artifact. The same server reuses a single keep-alive connection both under raw sequential HTTP requests and under the modern-era client, with or without the relay in the path. The mechanism is visible in the SDK. Legacy-era responses are SSE-framed, and the client closes each response stream as soon as the JSON-RPC reply event arrives; an HTTP client cannot return an undrained connection to its pool, so every exchange discards its connection. We found no client-side configuration that restores reuse in the legacy era. The modern era’s plain-JSON responses restore it automatically.² Our relay hides the cost of the extra connections; on a real 50 ms link they would add roughly another 100 ms of transport setup to the legacy MCP handshake, and more if TLS is in play.

Wire volume tells a consistent but secondary story. Full cold handshakes cost 1147+1522 B up/down for MCP (legacy) against 170+569 B for A2A. That difference matters at fleet scale but is invisible next to the round-trip effect at interactive timescales.

The phase decomposition (Figure 2) locates the cost. On loopback, most of every protocol’s handshake is client processing: SDK object construction, JSON parsing, validation. At 50 ms RTT the picture inverts. MCP (legacy) spends its time in session establishment and capability exchange (one round trip each, plus the readiness notification), MCP (modern) in capability exchange (discover + list), and A2A/ACP in discovery, each phase’s share tracking its round trips. There is no hidden compute cost that emerges at scale; the handshake is round trips nearly all the way down.

A sweep across emulated RTTs makes the mechanism direct (Figure 3). Median latency to task-ready rises linearly with RTT, with fitted slopes of 1.1 (A2A), 1.1 (ACP), 2.1 (MCP (modern)), and 4.2 (MCP (legacy)) milliseconds of handshake per millisecond of RTT ($N=20$ per cell). The sweep’s 50 ms points, measured in separate cells from those of Table 2, reproduce its medians within 2.4 ms. The single-fetch protocols track their one round trip almost exactly. The legacy MCP slope exceeds its three round trips because its two SSE-framed responses (`initialize` and `tools/list`)

²A related symptom of the same early-close pattern was reported upstream as `python-sdk` issue #2707 and closed for lack of reproduction; we filed a deterministic reproduction as issue #3281, <https://github.com/modelcontextprotocol/python-sdk/issues/3281>.




Figure 1: Application-layer round trips to task-ready (S1, S2) and to rejection (S3), by protocol variant. Warm values are what each protocol’s specification and official SDK allow a reconnecting client to skip; the pinned MCP variant exists only for warm reconnects and so appears only in the S2 group.

each cross the relay in two delayed chunks, while the `initialized` notification returns a bare 202: three delayed requests plus five delayed response chunks make eight one-way delays, or four RTT equivalents, which is where the fitted 4.2 lands (Section 4.5).

5.2 S2: warm repeat and amortization

Warm behavior separates the protocols more sharply than cold behavior (Table 2, Figure 1). MCP (legacy) can skip nothing. Its specification predates response caching, so a reconnecting client repeats all 3 round trips and pays 230 ms at 50 ms RTT, within noise of its cold handshake. MCP (modern) under the SDK’s default policy re-probes `server/discover` but serves the tool list from the SEP-2549 cache: 1 round trip, 63.5 ms. A client that additionally pins the counterpart’s protocol era reaches task-ready with zero wire exchanges, in 6.6 ms. That is the same zero-round-trip regime as A2A’s cached-card (2.0 ms) and ACP’s cached-manifest (6.6 ms) reconnects; the residual handful of milliseconds in each case is client-side setup, not wire exchange.

Figure 4 shows what this means over repeated interactions. At $k=100$ interactions with the same counterpart under 50 ms RTT, the per-interaction handshake cost amortizes to 230 ms for MCP (legacy), which is effectively no amortization at all, versus 64.1 ms for MCP (modern) (auto), and 7.8 ms, 2.6 ms, and 7.1 ms for pinned MCP, A2A, and ACP. For agent pairs that interact repeatedly, cacheability, not cold-start cost, is the design property that matters. For true strangers, the cold numbers stand.

5.3 S3: capability mismatch

Time-to-rejection mirrors time-to-ready, because in all three protocols the initiator discovers what the counterpart *cannot* do the same way it discovers what it can. A2A and ACP reject after their single metadata fetch (62.2 ms and 57.7 ms at 50 ms RTT), and MCP (modern) pays its discover-plus-list (117 ms). MCP (legacy) must complete its entire three-round-trip handshake, session included, before learning the counterpart is useless to it (231 ms). The wasted cost of a mismatched contact is thus the full handshake price, and it compounds for agents that shop across many counterparts. Document-centric discovery makes ruling out a counterpart as cheap as one GET; session-first designs charge the session price merely to browse. The ACP medians invert slightly across these two scenarios (57.7 ms to rejection versus 62.1 ms to cold task-ready). The per-request logs attribute the entire difference to client-side setup before the request leaves the client (wire-level GET times agree to within a millisecond), which drifts by a few milliseconds between cells over a long benchmarking session. Cross-cell differences below roughly 5 ms should be read as this noise floor, not as protocol




Figure 2: S1 cold-start handshake latency decomposed by phase (median stacked; whisker to p95 of the total), on loopback and under the emulated 50 ms RTT. Note the differing y-scales: loopback cost is client processing; at 50 ms RTT the wire phases dominate in proportion to their round trips.

Table 3: Behavior under degraded conditions ($N=50$ per cell, localhost). S4a reports time to task-ready with a 2 s per-response server delay; S4b/S4c report time to failure. Retr. = client retries observed (duplicate requests). For S4c, ACP is untestable: its manifest carries no protocol-version field.

Scenario	Protocol	Outcome	med. time (ms)	Retr.	Error surfaced
S4a slow (2 s/resp.)	MCP (legacy)	ready	6026	0	none
	MCP (modern)	ready	4019	0	none
	A2A	ready	2011	0	none
	ACP	ready	2011	0	none
S4b truncated metadata	MCP (legacy)	error	3.5	0	MCPError
	MCP (modern)	error	4.4	0	MCPError
	A2A	error	3.0	0	AgentCardResolutionError
	ACP	error	2.9	0	JSONDecodeError
S4c version mismatch	MCP (legacy)	error	3.7	0	RuntimeError
	MCP (modern)	error	4.2	0	RuntimeError
	A2A	ready	3.0	0	none (version deferred)

effects. Note that this is a systematic cross-cell term, not sampling error: the bootstrap intervals within each cell are an order of magnitude tighter (Section 4), which is why the two cells can be individually well resolved and still not differ meaningfully from each other.

5.4 S4: degraded conditions

Table 3 summarizes the degraded scenarios.

S4a: slow counterpart. With 2 s injected per response, time to task-ready is a clean multiple of round-trip count: 6026 ms for MCP (legacy) (three delayed responses), 4019 ms for MCP (modern), and about 2 s for A2A and ACP (2011 ms, 2011 ms). No SDK timed out or retried; the delay simply multiplies. A slow counterpart is *most* expensive to the protocol with the most round trips, and the initiator has no way to discover the slowness before paying the full handshake.

S4b: malformed metadata. All four variants fail fast, with no retry storms and no hangs. Every client detected truncated JSON on first parse, in milliseconds. What differs is error quality. A2A’s SDK raises a typed `AgentCardResolutionError` naming the URL. MCP’s raises a protocol-level `MCPError` wrapping the JSON validation failure (the modern client, having received a malformed discover response, falls back to `initialize` per its negotiation rule and fails cleanly




Figure 3: Cold-start latency to task-ready versus emulated RTT (S1 medians, $N=20$ per cell). Slopes track the number of delayed exchanges per handshake. Right-edge leader lines connect curves to labels and are not data.

when that response is also malformed: two requests, no blind retries). ACP’s client surfaces a bare `JSONDecodeError` with a character offset and no indication of which endpoint or agent was involved. All three abort rather than degrade, which is the right behavior; only ACP leaves the caller to reconstruct the context.

S4c: version mismatch. Here the protocols genuinely diverge. Both MCP eras refuse a counterpart that speaks only the fictitious version 2099-01-01, in 3.7–4.2 ms, with an explicit error naming the unsupported version; the modern probe’s structured `-32022` error even carries the server’s supported-version list. A2A’s client behaved differently. It fetched a card advertising protocol version 99.0.0, selected the JSON-RPC binding, and completed echo tasks in normal handshake time (3.0 ms) without surfacing any warning. This is spec-conformant client behavior, not an SDK bug. The normative schema marks the interface’s `protocolVersion` as descriptive (“the version of the A2A protocol this interface exposes”), and A2A’s mandatory version negotiation instead happens per request, via an `A2A-Version` header that servers must reject when unsupported [1]. On the wire we observed exactly that split. The client’s task request carried `a2a-version: 1.0`, its own version, independent of the card’s claim, and our genuinely v1.0 counterpart accepted it. The consequence for handshake design stands: discovery-time version claims go unvalidated, so a real incompatibility surfaces not during the handshake but at the first task message, as a server-side rejection, and when the wire versions happen to agree, as here, it never surfaces at all. ACP has no version signal to check. Deferred version checking is convenient exactly until an incompatibility is real, at which point what could have been a cheap, legible handshake failure becomes a task-time failure attributed to the wrong layer.

5.5 Context cost of capability metadata

When the initiating agent is an LLM, the counterpart’s metadata is not merely transferred, it is *read*: placed in a context window to select a tool, check a skill, or validate a counterpart. The scarce




Figure 4: Amortized handshake cost per interaction over k interactions with the same counterpart (cold $+(k-1) \times$ warm, over k), under the emulated 50 ms RTT. Protocols with zero-round-trip warm paths amortize to a few milliseconds of client-side setup; MCP (legacy) pays its full handshake forever. Right-edge leader lines connect curves to labels and are not data.

resource is therefore context, and the relevant question is what fraction of the window the handshake consumes before any task content arrives.

For our one-capability counterpart the answer is trivial (Table 4): 112–242 tokens, a rounding error in any modern context window. That number is a protocol floor and should not be quoted as a cost. It matters only as an intercept.

The slope is the finding (Table 5). Each additional capability costs 24 tokens in an A2A card, 78 in an MCP tool listing, and 107 in an ACP manifest. At 200 capabilities, a plausible size for a mature tool server, the handshake alone delivers 15,618 tokens for MCP and 21,405 for ACP, which is 11% of a 200k-token window spent before the agent has read its own task. Three points follow.

First, the per-capability costs differ by more than a factor of four, and the ordering is the reverse of the one-capability ranking. A2A is cheapest per skill because a card entry is a name, a description, and tags; MCP is dearer because a tool entry carries full input and output JSON schemas; ACP is dearest because each agent entry repeats a sixteen-field metadata block that is null in our counterpart and would be populated in a real one. A protocol’s floor tells you little about what it costs at deployment scale.

Second, this is the one handshake cost that round-trip reduction does not touch. It is carried by the *listing* step, which every protocol performs and which the 2026 MCP revision made cacheable but no cheaper. Indeed the modern era spends tokens to save round trips: its richer discovery envelope makes it the most token-expensive handshake at $n=1$ (242 versus 181), a trade that pays for itself only because the result is cacheable.

Table 4: Capability metadata an LLM-driven initiator ingests during one cold handshake with our one-capability counterpart (exact wire bodies; o200k_base tokens, a lower bound; see Section 4).

Protocol	metadata bytes	tokens	USD / 10^6 handshakes
MCP (legacy)	726	181	\$543
MCP (modern)	953	242	\$726
A2A	443	113	\$339
ACP	507	112	\$336

Table 5: Handshake metadata in tokens as the counterpart’s capability inventory grows, and the marginal cost of one capability. Built by replicating the captured wire entry n times inside the byte-identical envelope, so the documents are exactly what a counterpart with n such capabilities would serve.

Protocol	$n=1$	$n=10$	$n=50$	$n=200$	per cap.
MCP	96	798	3,918	15,618	78
A2A	115	331	1,291	4,891	24
ACP	112	1,075	5,355	21,405	107

Third, cacheability is what bounds the damage. The same SEP-2549/cached-card machinery that eliminates warm round trips keeps already-ingested metadata valid, so an agent that reconnects pays the context cost once rather than per interaction. For an agent that shops across many counterparts, however, there is no reuse to be had, and inventory size becomes the binding constraint on how many counterparts it can consider at all.

5.6 Transport-layer baseline

For calibration: on the same loopback, a bare TCP connect completes in 0.05 ms (median) and a full TLS 1.3 handshake in 1.88 ms of compute. On a real network each adds one round trip [19]. Against that two-round-trip transport-plus-cryptographic stack, A2A and ACP add a single application-layer round trip, half the cost of the entire stack beneath them. A legacy MCP handshake adds three, exceeding the combined transport and cryptographic setup it rides on, and that is before any authentication, which none of our configurations exercised. The application layer has, in effect, re-created the pre-TLS-1.3 handshake stack one level up, and it is now re-learning the same remedies (fewer round trips, resumable state) that transport designers adopted a decade ago [11, 12].

6 Discussion

Where the overhead concentrates. Our measurements are consistent across scenarios. Handshake cost is round trips first, connections second, bytes a distant third, and compute negligible. Every design question about agent handshakes should therefore be asked as “how many round trips does this add, and can it be cached?” The protocols’ own trajectories agree. MCP’s 2026 revision is precisely a round-trip-reduction release: it replaced a three-message, session-first initialization with a one-probe, cacheable discovery, and our measurements quantify the effect at 47% off the cold handshake at 50 ms RTT and, with pinning, free warm reconnects. That the youngest protocol in our study already shipped this change is itself evidence that handshake overhead is a live design pressure, not a hypothetical.

Document-centric versus session-first discovery. A2A and ACP demonstrate the value of putting the entire capability surface in one cacheable document at a well-known location: cold task-ready in one round trip, mismatch rejection at the same price, warm reconnects free. This design has two costs. One is currency, since a cached card can be stale and neither protocol gives the initiator a cheap way to validate freshness short of re-fetching. The other is expressiveness, since a static document cannot answer negotiated or per-principal capability questions. MCP’s modern era lands in between. It keeps a negotiated, session-capable protocol but makes the negotiation’s *results* cacheable with

explicit server-controlled TTLs, which preserves freshness semantics the static documents lack.

Failure legibility is a protocol property. The degraded scenarios show that what a protocol checks is as consequential as what it exchanges. MCP’s explicit version negotiation bought it clean, fast, well-attributed failures in every degraded cell. A2A, by design, defers version enforcement to per-request headers checked by the server, so its discovery-time version claims pass unexamined. ACP’s missing version field forecloses the check entirely. For autonomous agents, which cannot eyeball a stack trace, the difference between a structured “unsupported version, here is what I support” and a counterpart whose incompatibility surfaces only at task time is the difference between a recoverable negotiation and a debugging session.

Implications for agent-system designers. Three follow directly from the data. First, treat counterpart metadata as a cacheable asset with a lifecycle: the measured gap between MCP (legacy) and everything else is mostly the gap between re-deriving and reusing handshake state. Second, budget for round trips at discovery time when agents shop across many counterparts. S3 shows rejection costs a full handshake, so a design that can rule counterparts out from a one-GET document scales shopping linearly in counterparts, while session-first designs scale it in handshakes. Third, keep capability inventories lean if an LLM will read them. Section 5.5 shows the per-capability slope, not the floor, is what binds: listing is the step that grows with inventory, and it is the one handshake cost that caching defers but round-trip reduction never touches. We deliberately stop at implications. Whether a future protocol should, for example, merge A2A-style cards with MCP-style negotiated sessions is a design question our measurements inform but do not settle.

7 Limitations and Future Work

Beyond the validity threats of Section 4.5, three boundaries mark natural extensions. We measured anonymous configurations; authenticated handshakes (OAuth flows in MCP, credentialed A2A cards) add round trips and token exchanges that deserve the same decomposition. We measured one counterpart per protocol; sweeping capability-inventory size would turn our token floors into curves. And we excluded live natural-language negotiation in the style of Agora [15], where two LLM-backed agents converge on a shared task format before first use. By Agora’s own accounting such negotiation costs thousands of tokens per pair, orders of magnitude above every structured handshake we measured, and benchmarking it under our phase decomposition, including its amortization across subsequent structured interactions, is the most interesting open cell in our matrix. Finally, our S2 results suggest that hybrid designs (cacheable discovery documents carrying negotiated session hints, or 0-RTT-style resumption tokens for agent sessions) could combine the strengths we measured separately. Evaluating such designs is future work, not a proposal we make here.

8 Conclusion

To our knowledge, this is the first empirical measurement of the agent handshake decomposed into its phases (discovery, capability exchange, session, readiness), across the major AI agent communication protocols and across cold, warm, mismatch, and degraded scenarios, with an open harness and raw data. The tax of saying hello is real and structured. It is paid in round trips, multiplied by network latency and counterpart slowness, waived by cacheability, and, in one protocol’s own 2026 revision, already being engineered down. Agent-to-agent interaction is young enough that its connection economics are still being designed. We hope a shared, reproducible way to measure the cost of hello helps that design conversation stay empirical.

References

- [1] A2A Project, Linux Foundation. Agent2Agent (A2A) protocol specification, v1.0.0. <https://a2a-protocol.org/v1.0.0/specification>, 2026. Released 2026-03-12; accessed 2026-08-10.

- [2] Zeynab Anbiaee, Mahdi Rabbani, Mansur Mirani, Gunjan Piya, Igor Opushnyev, Ali Ghorbani, and Sajjad Dadkhah. Security threat modeling for emerging AI-agent protocols: A comparative analysis of MCP, A2A, Agora, and ANP, 2026. URL <https://arxiv.org/abs/2602.11327>.
- [3] Anthropic. Introducing the Model Context Protocol. <https://www.anthropic.com/news/model-context-protocol>, 2024. Published 2024-11-25.
- [4] Ivan Dobrovolskyi. Empirical comparison of agent communication protocols for task orchestration, 2026. URL <https://arxiv.org/abs/2603.22823>.
- [5] Hongyi Du, Jiaqi Su, Jisen Li, Lijie Ding, Yingxuan Yang, Peixuan Han, Xiangru Tang, Kunlun Zhu, and Jiaxuan You. ProtocolBench: Which LLM multi-agent protocol to choose? In *Proceedings of the 43rd International Conference on Machine Learning (ICML)*, 2026. URL <https://arxiv.org/abs/2510.17149>. To appear; arXiv:2510.17149.
- [6] Abul Ehtesham, Aditi Singh, Gaurav Kumar Gupta, and Saket Kumar. A survey of agent interoperability protocols: Model Context Protocol (MCP), Agent Communication Protocol (ACP), Agent-to-Agent Protocol (A2A), and Agent Network Protocol (ANP), 2025. URL <https://arxiv.org/abs/2505.02279>.
- [7] Envoy Project. xDS REST and gRPC protocol. https://www.envoyproxy.io/docs/envoy/latest/api-docs/xds_protocol, 2026. Accessed 2026-08-11.
- [8] Foundation for Intelligent Physical Agents. FIPA ACL message structure specification. Technical Report SC00061G, FIPA, 2002. URL <http://www.fipa.org/specs/fipa00061/SC00061G.html>.
- [9] Xinyi Hou, Yanjie Zhao, Shenao Wang, and Haoyu Wang. Model Context Protocol (MCP): Landscape, security threats, and future research directions, 2025. URL <https://arxiv.org/abs/2503.23278>.
- [10] IBM BeeAI project. Agent Communication Protocol specification. <https://agentcommunicationprotocol.dev/>, 2025. Version 1.0; repository archived 2025-08-27; accessed 2026-08-10.
- [11] Jana Iyengar and Martin Thomson. QUIC: A UDP-based multiplexed and secure transport. Technical Report RFC 9000, IETF, 2021.
- [12] Hyunwoo Lee, Doowon Kim, and Yonghwi Kwon. TLS 1.3 in practice: How TLS 1.3 contributes to the internet. In *Proceedings of the Web Conference (WWW)*, 2021. doi: 10.1145/3442381.3450057.
- [13] LF AI & Data Foundation. ACP joins forces with A2A. <https://lfaidata.foundation/communityblog/2025/08/29/acp-joins-forces-with-a2a-under-the-linux-foundations-lf-ai-data/>, 2025. Published 2025-08-29.
- [14] Yedidel Louck, Ariel Stulman, and Amit Dvir. Security analysis of agentic AI communication protocols: A comparative evaluation, 2025. URL <https://arxiv.org/abs/2511.03841>.
- [15] Samuele Marro, Emanuele La Malfa, Jesse Wright, Guohao Li, Nigel Shadbolt, Michael Wooldridge, and Philip Torr. A scalable communication protocol for networks of large language models, 2024. URL <https://arxiv.org/abs/2410.11905>.
- [16] Model Context Protocol project. Model Context Protocol specification, revision 2025-11-25. <https://modelcontextprotocol.io/specification/2025-11-25>, 2025. Accessed 2026-08-10.

- [17] Model Context Protocol project. Model Context Protocol specification, revision 2026-07-28. <https://modelcontextprotocol.io/specification/2026-07-28>, 2026. Accessed 2026-08-10.
- [18] David Naylor, Alessandro Finamore, Ilias Leontiadis, Yan Grunenberger, Marco Mellia, Maurizio Munafò, Konstantina Papagiannaki, and Peter Steenkiste. The cost of the “s” in HTTPS. In *Proceedings of the 10th ACM International Conference on Emerging Networking Experiments and Technologies (CoNEXT)*, pages 133–140, 2014. doi: 10.1145/2674005.2674991.
- [19] Eric Rescorla. The Transport Layer Security (TLS) protocol version 1.3. Technical Report RFC 8446, IETF, 2018.
- [20] Reid G. Smith. The Contract Net Protocol: High-level communication and control in a distributed problem solver. *IEEE Transactions on Computers*, C-29(12):1104–1113, 1980. doi: 10.1109/TC.1980.1675516.
- [21] Martin Thomson and Cory Benfield. HTTP/2. Technical Report RFC 9113, IETF, 2022.
- [22] Yingxuan Yang, Huacan Chai, Yuanyi Song, Siyuan Qi, Muning Wen, Ning Li, Junwei Liao, Haoyi Hu, Jianghao Lin, Gaowei Chang, Weiwen Liu, Ying Wen, Yong Yu, and Weinan Zhang. A survey of AI agent protocols, 2025. URL <https://arxiv.org/abs/2504.16736>.
- [23] Xiangfeng Zhu, Guozhen She, Bowen Xue, Yu Zhang, Yongsu Zhang, Xuan Kelvin Zou, Xiongchun Duan, Peng He, Arvind Krishnamurthy, Matthew Lentz, Danyang Zhuo, and Ratul Mahajan. Dissecting overheads of service mesh sidecars. In *Proceedings of the 2023 ACM Symposium on Cloud Computing (SoCC)*, 2023. doi: 10.1145/3620678.3624652. Extended version: arXiv:2207.00592.

A Reproducibility

The repository at

<https://github.com/pvarshh/handshake-benchmark>

contains the harness (/harness), the minimal counterpart agents (/agents), raw per-run JSONL for every cell (/results/raw), and the scripts that generate every figure, table, and in-text number (/analysis). Three commands reproduce the raw data, and one regenerates every figure, table, and number in this paper from it:

```
python -m harness.run
    --scenario all --n 50 --network both
python -m harness.baseline
python -m harness.rtt_sweep
python -m analysis.run_all
```

The first takes roughly 25 minutes on a laptop; the others are quick. Version pins: Python 3.12.13, mcp 2.0.0, a2a-sdk 1.1.2, acp-sdk 1.0.3, httpx 0.28.1, uvicorn 0.34.3, fastapi 0.115.14, starlette 0.46.2. The MCP counterpart is served over streamable HTTP with SEP-2549 cache hints (ttlMs = 3,600,000, i.e. one hour, public scope) on server/discover and tools/list; the A2A counterpart serves a v1.0 Agent Card with one skill over the JSON-RPC binding; the ACP counterpart serves one agent manifest. Degraded variants are driven by environment flags documented in the repository README.